Group Project
Data Science and Machine Learning

Authors:
Jonas Brandenburg if20b243, Stephan Düx if20b245, Violeta Garcia Espin if20b091, Rawan Mousa if20b129

1. Data Acquisition

The **Behavioral Risk Factor Surveillance System** (BRFSS) is the nation's premier system of health-related telephone surveys that collect state data about U.S. residents regarding their health-related risk behaviors, chronic health conditions, and use of preventive services.

- 400,000 adult interviews each year
- Largest continuously conducted health survey in the world
- 336 columns
- ~600mb of data per year
- Datasets from 1986 - 2020 available

https://www.cdc.gov/brfss/annual_data/annual_data.htm

https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system

2. Data Exploration

441.456 Samples
330 Features

2. Data Exploration: Feature Selection

2. Data Preprocessing: Feature Selection

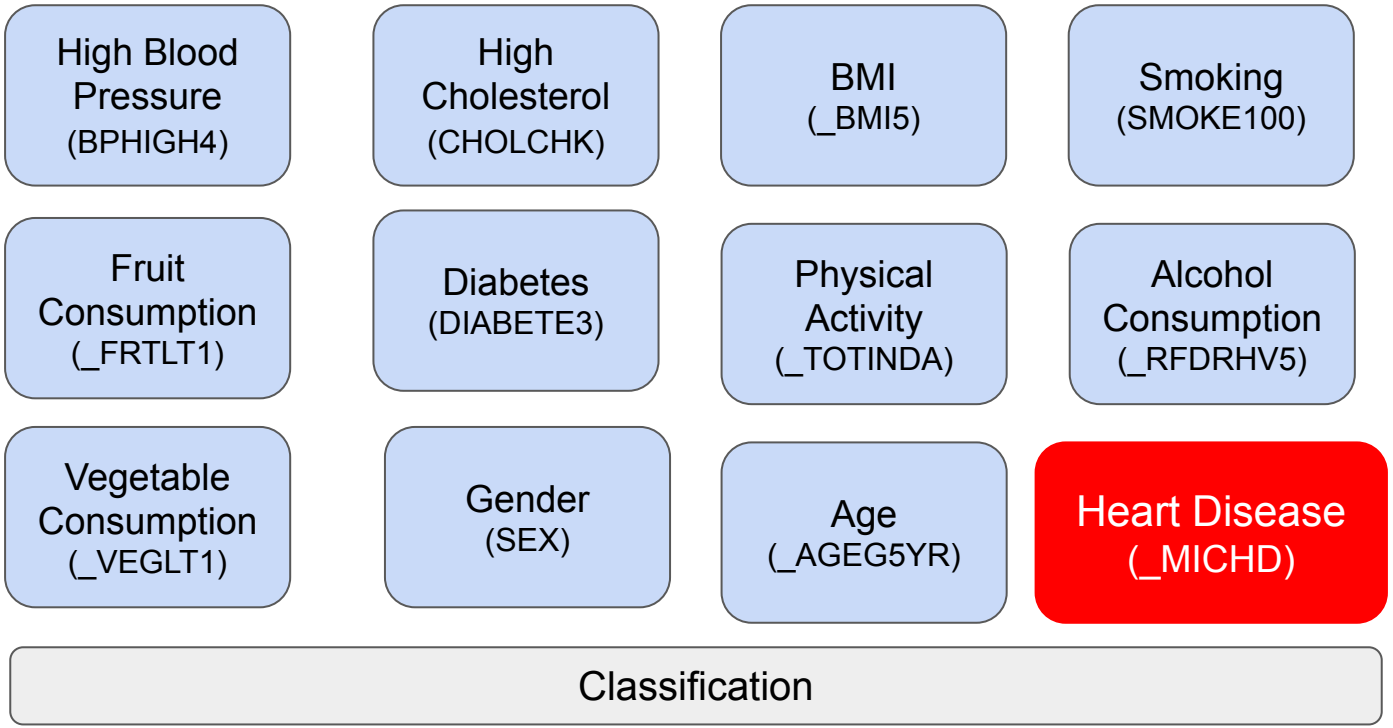

https://www.cdc.gov/brfss/annual_data/2015/pdf/codebook15_llcp.pdf

3. Data Processing: Label Encoding

Feature: HighBP (_RFHYPE5)

Adults who have been told they have high blood pressure by a doctor, nurse, or other health professional.

Labels:

- 1: No
- 2: Yes
- 9: Don't know

For logic purposes, label 1 will be substitute by 0 and label 2 by 1.

In 30

```
# HighBP
df['HighBP'] = df['HighBP'].replace({9:np.nan})
df['HighBP'] = df['HighBP'].replace({1:0})
df['HighBP'] = df['HighBP'].replace({2:1})
```

3. Data Preprocessing: Label Encoding

One Hot Encoding would have been indicated.

Feature Diabetes (DIABETE3)

People that have been told to suffer from diabetes.

Labels:

- 1: yes
- 2: yes but only during pregnancy
- 3: no
- 4: pre-diabetes or borderline diabetes.
- 7: Don't know
- 9: refuse to answer

We will just consider yes(1) or no (3) answers. To simplify, the rest of the options will be deleted.

In 33

```
# Diabetes
df['Diabetes'] = df['Diabetes'].replace({2:np.nan})
df['Diabetes'] = df['Diabetes'].replace({4:np.nan})
df['Diabetes'] = df['Diabetes'].replace({7:np.nan})
df['Diabetes'] = df['Diabetes'].replace({9:np.nan})
df['Diabetes'] = df['Diabetes'].replace({3:0})
```

4. Data Preprocessing: Handling NaN Values

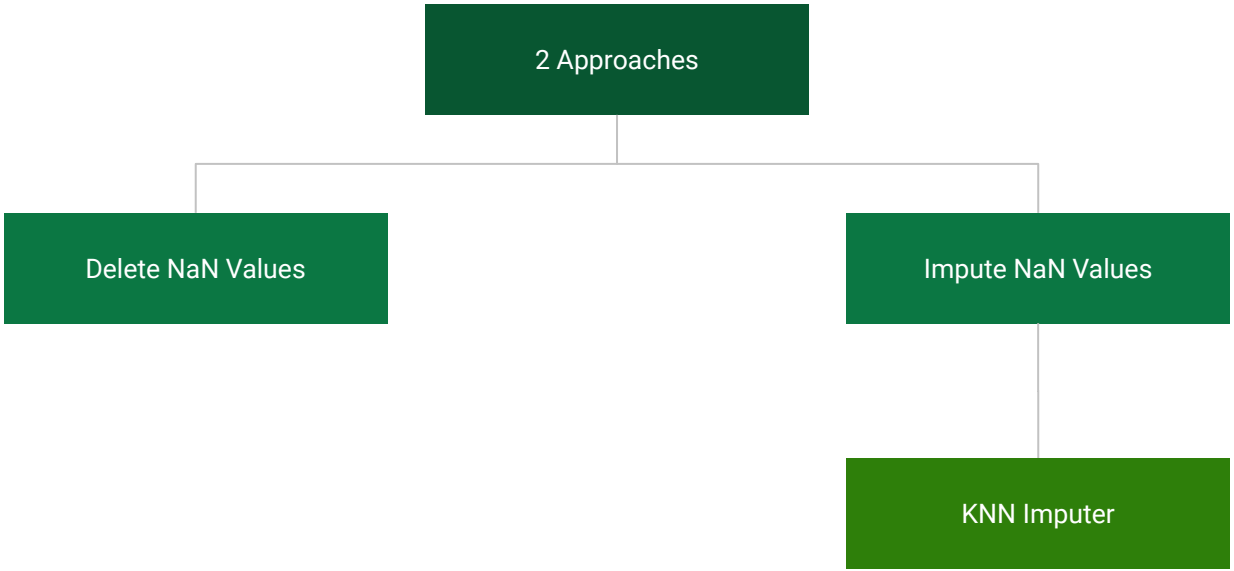

5. Data Preprocessing: Examination of individual features

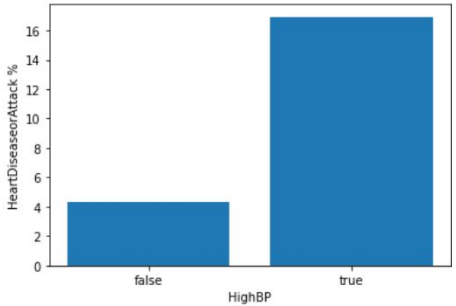

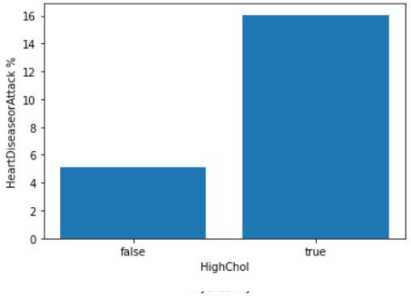

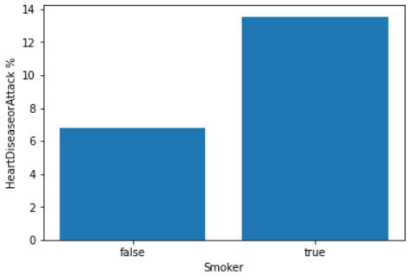

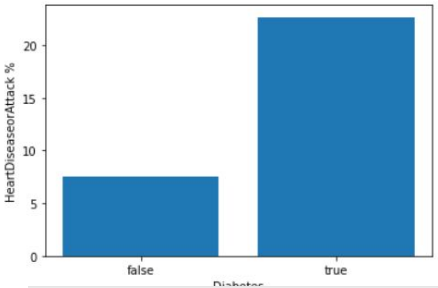

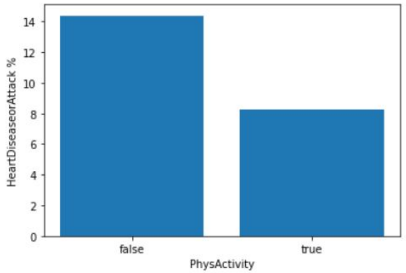

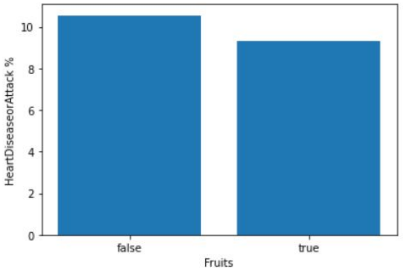

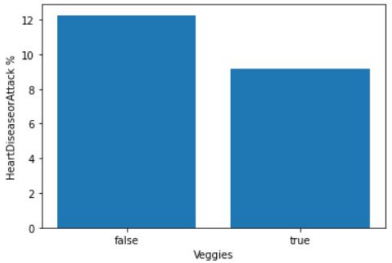

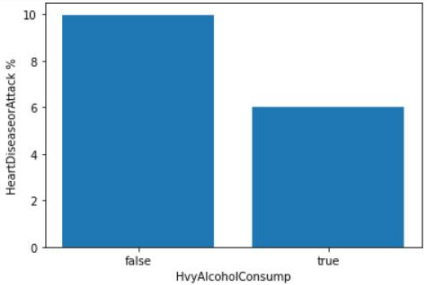

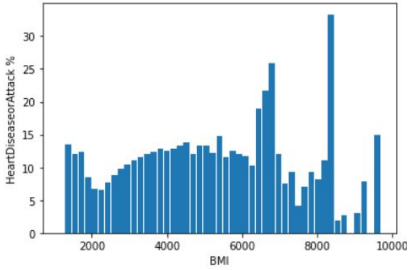

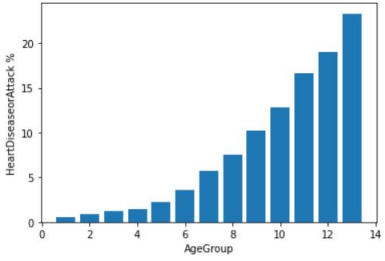

6. Data Preprocessing: Balanced data

	index	HeartDiseaseorAttack
0	0.0	264137
1	1.0	28608

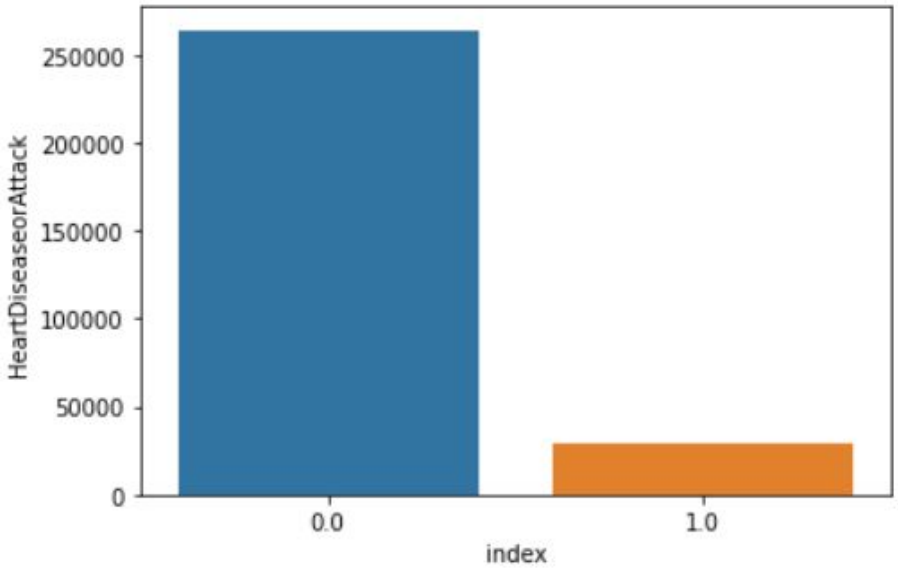

7. Model: KNN - Hyperparameter

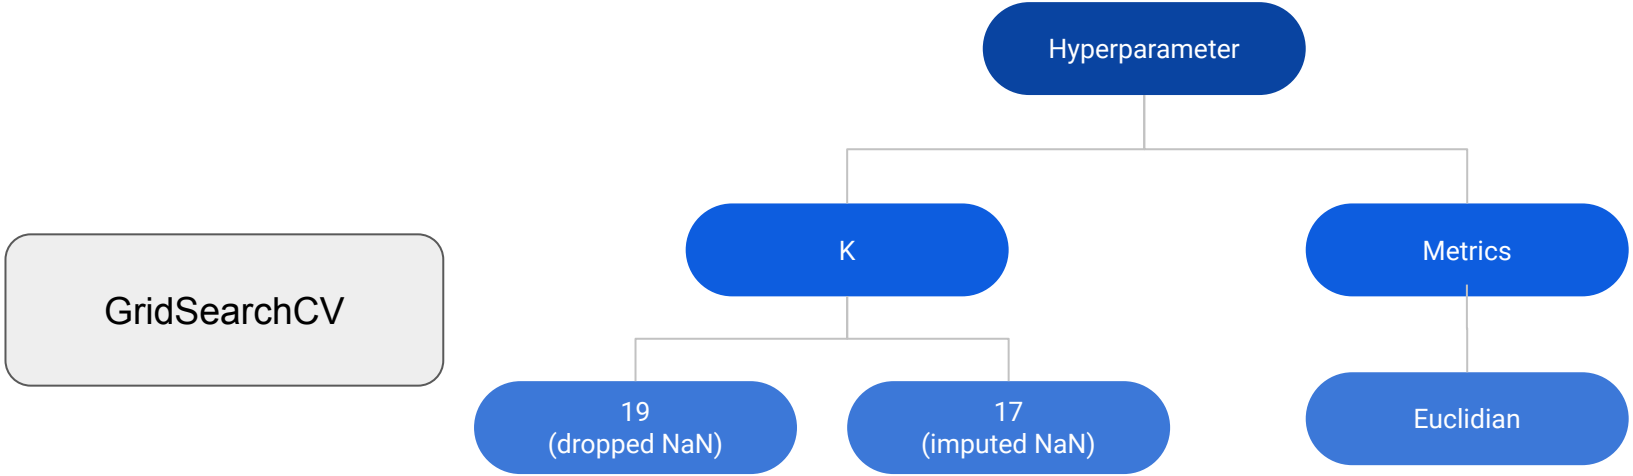


```
[ ]  from sklearn.metrics import DistanceMetric
     params = {'n_neighbors' : k, 'metric': ['manhattan','euclidean', 'minkowski']}
```

```
[ ]  random_search = GridSearchCV(knn, params, cv=5)
     random_search.fit(X_train, y_train)
```

7. Model: KNN

```
[37] accuracy_list =[]
     for i in range(100):
       X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size = 0.2)
       knn = KNeighborsClassifier(n_neighbors=best_k, metric = metric)
       knn.fit(X_train, y_train)
       prediction = knn.predict(X_test)

       balanced_accuracy_score(y_test, prediction)
       accuracy_list.append(balanced_accuracy_score(y_test, prediction))
     print("Balanced Accuracy Score is: ", np.mean(accuracy_list))
```

7. Model KNN - Performance

Accuracy = 0.72 for dataset with dropped NaN values

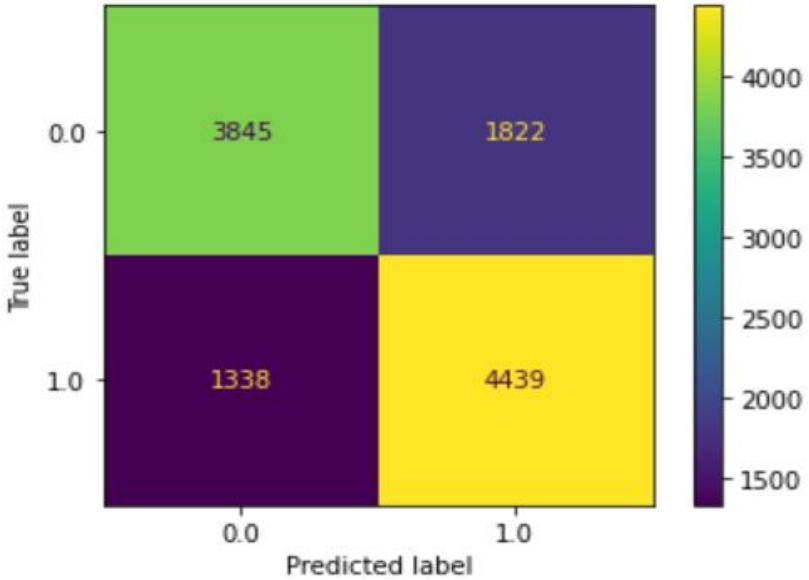

Accuracy = 0.73 for dataset with imputed NaN values

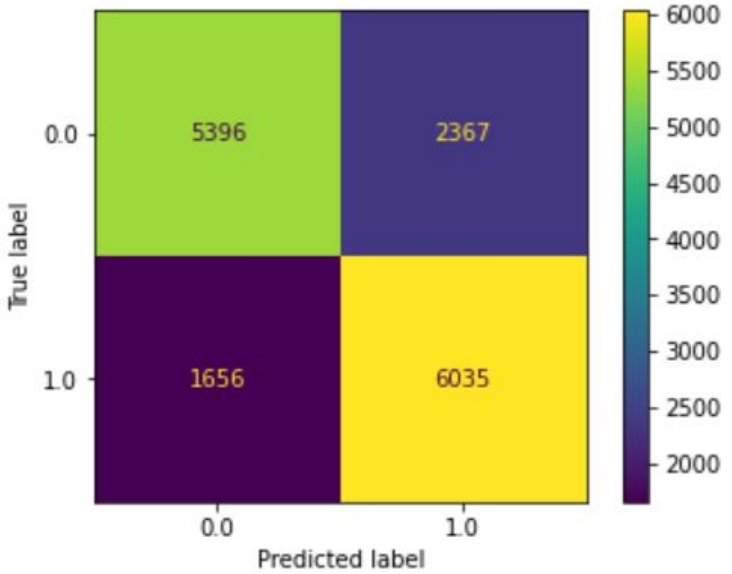

7. Model KNN - Performance (k vs. Accuracy)

dropped NaN values

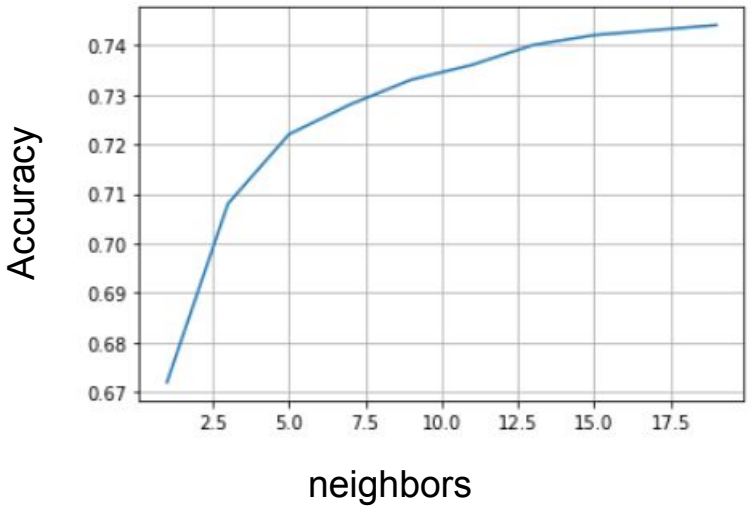

Imputed NaN values

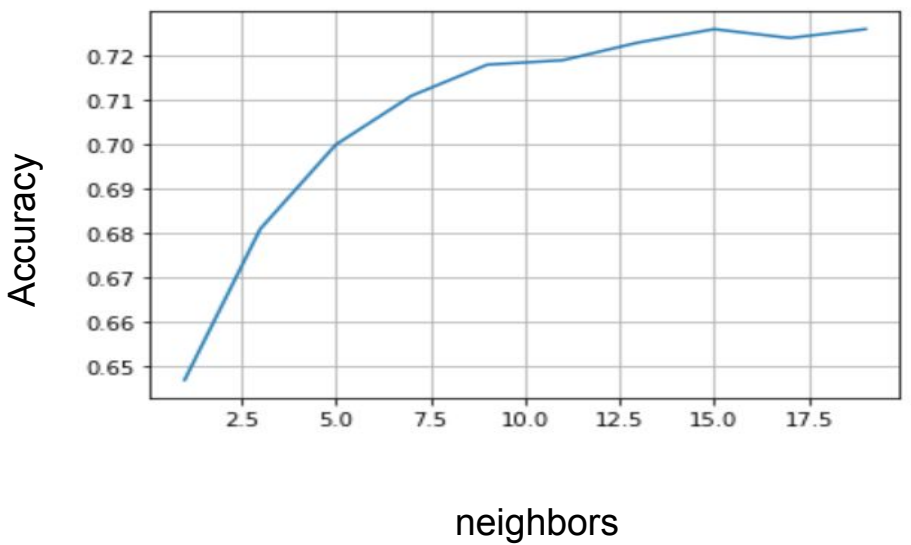

8. Model: Decision Trees - Gini vs. Entropy

```
[151] gini_scores  = []
      entropy_scores = []

      for i in range (0, 100):
        X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size = 0.2)
        gini_model = DecisionTreeClassifier(criterion='gini', random_state=1)
        entropy_model = DecisionTreeClassifier(criterion='entropy', random_state=1)

        gini_model.fit(X_train, y_train)
        entropy_model.fit(X_train, y_train)

        gini_predictions = gini_model.predict(X_test)
        entropy_prediction = entropy_model.predict(X_test)

        gini_scores.append(balanced_accuracy_score(y_test, gini_predictions))
        entropy_scores.append(balanced_accuracy_score(y_test, entropy_prediction))

      avg_scores_gini = sum(gini_scores) / len(gini_scores)
      avg_scores_entropy = sum(entropy_scores) / len(entropy_scores)

      print(f"Average accuracy score for gini {avg_scores_gini}")
      print(f"Average accuracy score for entropy {avg_scores_entropy}")
```

Gini	Entropy	
0.642	0.647	Dropped NaN
0.673	0.674	Imputed NaN

Entropy might perform slightly better but computationally more demanding.

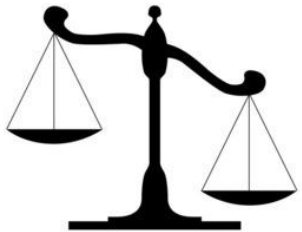

8. Model: Decision Trees

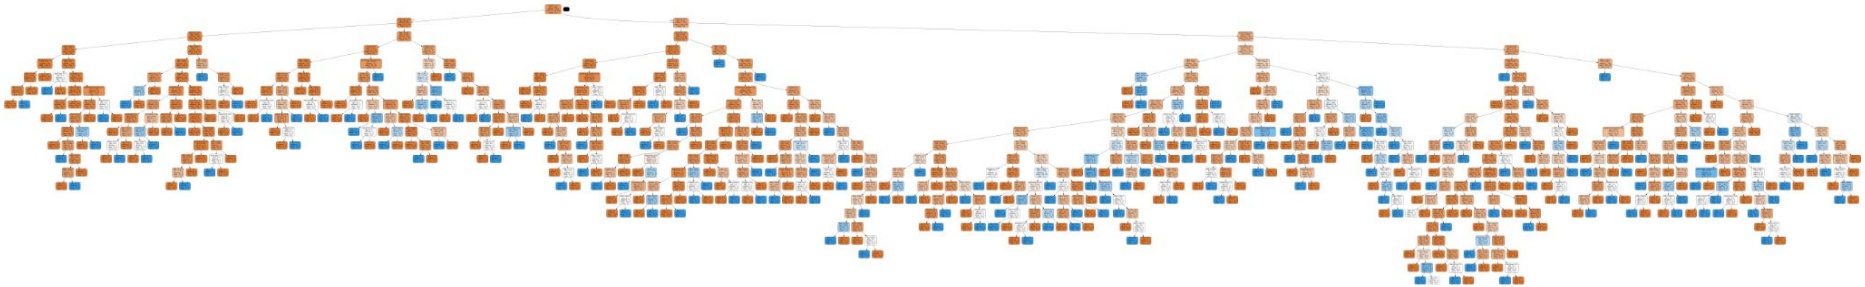

Source:
https://www.123rf.com/photo_138703028_vector-illustration-of-a-boy-pruning-bush-with-shears-isolated-on-white-background-cute-kid-doing-ga.html?vti=oenrrzu1j54jpx7ld5-1-6

8. Model: Decision Trees - Pruning

Pre-Pruning:

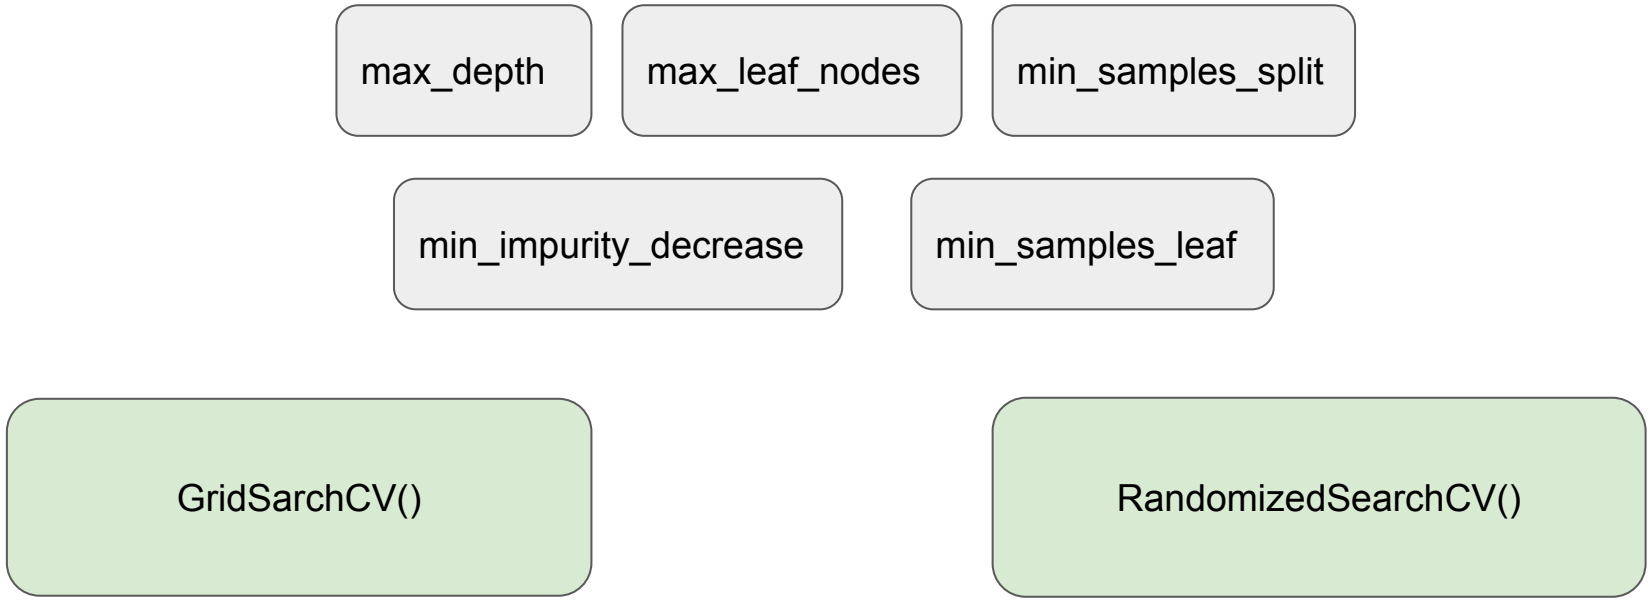

8. Model: Decision Trees - Pruning

Post-Pruning: ccp_alpha

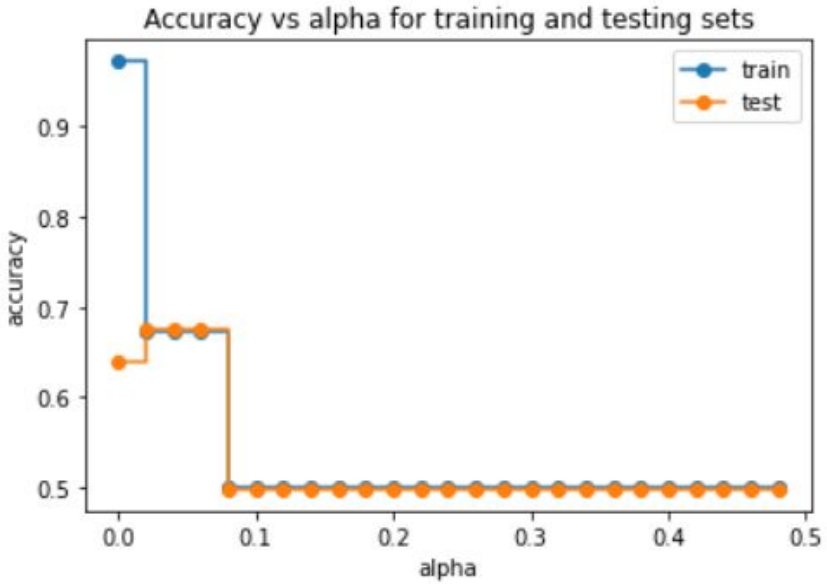

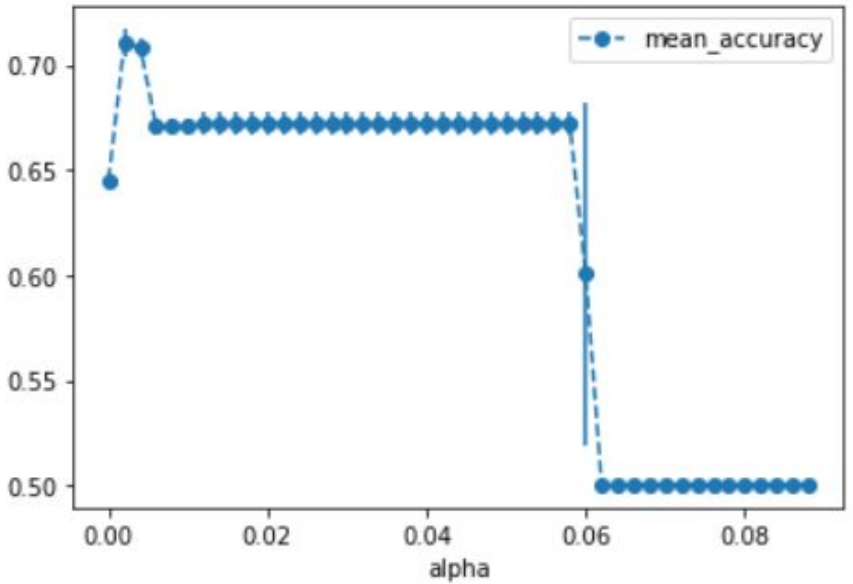

8. Model: Decision Trees - Optimized Tree

Hyperparameters:
ccp=0.002, criterion=gini

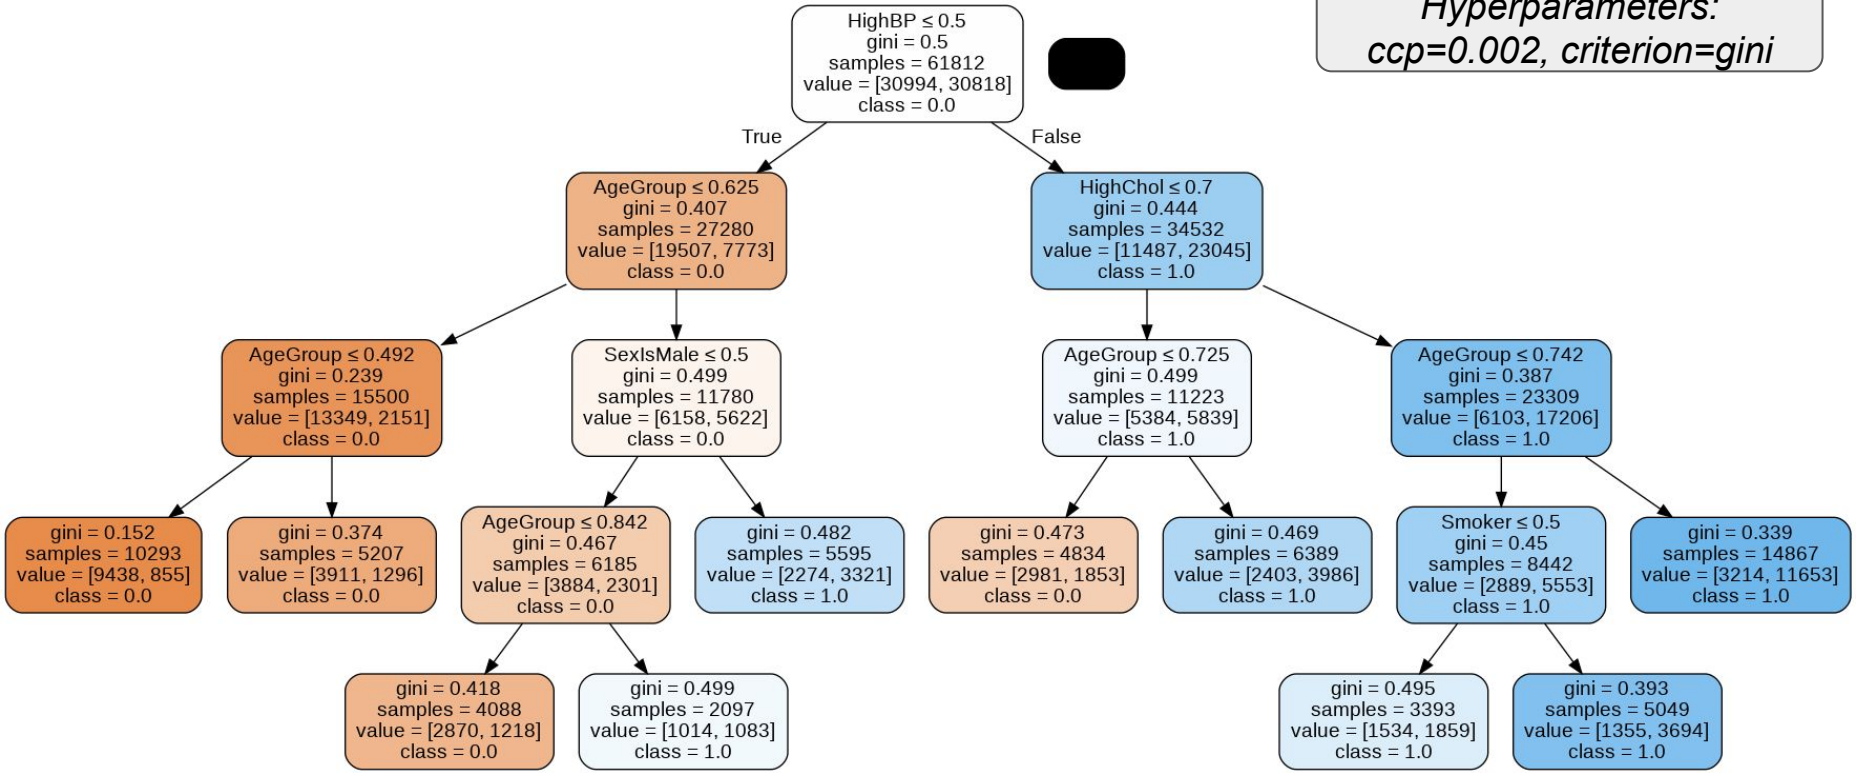

8. Model: Decision Trees - Performance

Accuracy = 0.708 for dataset with dropped NaN values

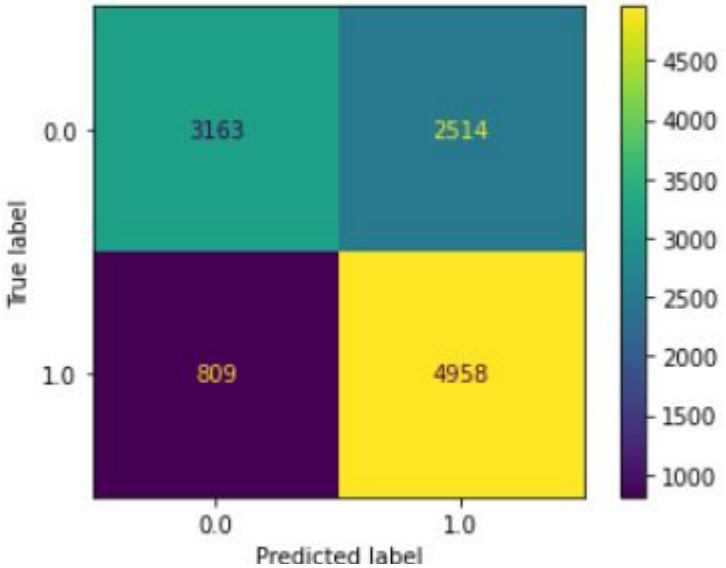

Accuracy = 0.728 for dataset with imputed NaN values

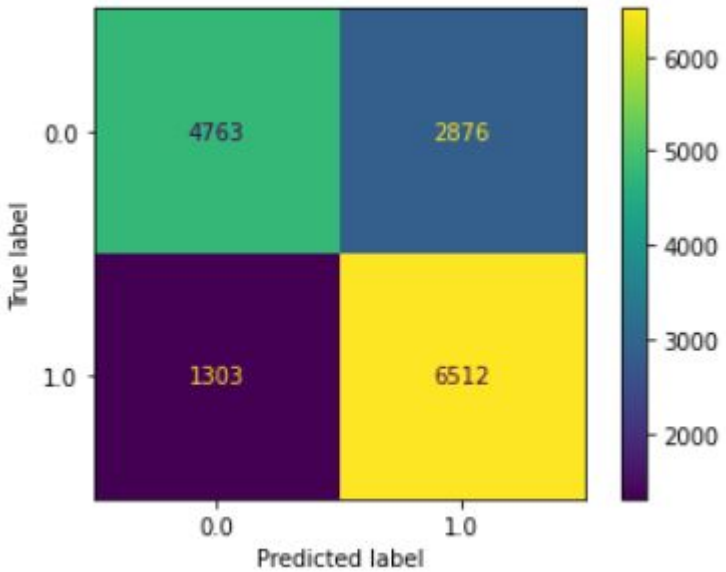

9. Model: MLP - Overview

Hyper Parameters

Type	Values		
KNN Imputation	NONE	KNN	
Activation function	sigmoid	relu	
ANN Arch: hidden layers	1	2	3
ANN Arch: hidden layers nodes	10	20	50
solver	adam	sgd	

2 x 2 x 3 x 3 x 2 = 72 combinations

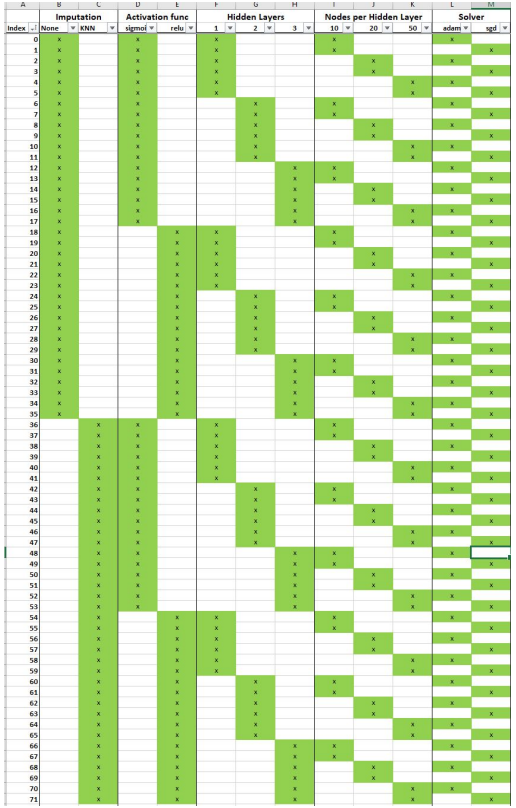

9. Model: MLP - Activation Function

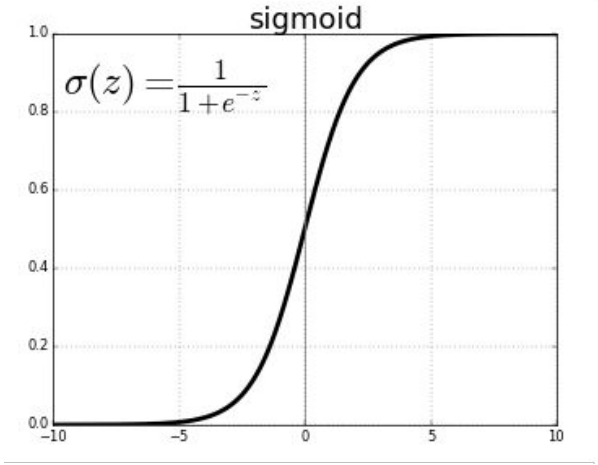

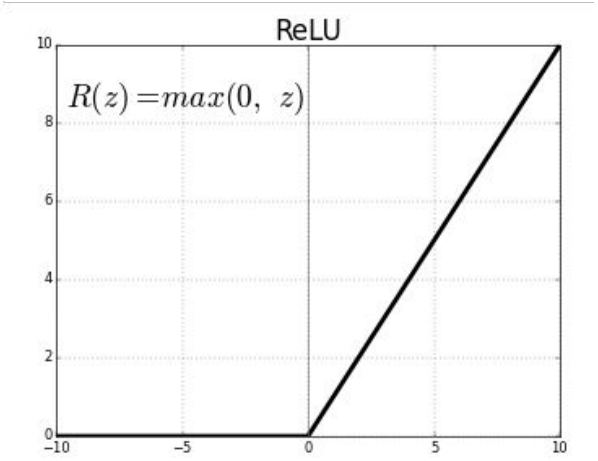

Activation func		Results	
sigmoi	relu	Balanced acc	Clears result ga
	x	0.752983544	TRUE
	x	0.752917655	TRUE
	x	0.752209507	TRUE
	x	0.752151507	TRUE
	x	0.751377622	TRUE
	x	0.751052197	TRUE
	x	0.751049684	TRUE
	x	0.751041243	TRUE
	x	0.750471004	TRUE
x		0.750403557	TRUE
x		0.750332744	TRUE
	x	0.750138846	TRUE
x		0.749946103	TRUE
x		0.749814174	TRUE
x		0.749810506	TRUE
x		0.749680387	TRUE
	x	0.749498346	TRUE
	x	0.749498246	TRUE
	x	0.749484076	TRUE
x		0.749174228	TRUE
	x	0.749172419	TRUE
x		0.749171112	TRUE
	x	0.749039033	TRUE
	x	0.74885026	TRUE
x		0.748849054	TRUE
x		0.748721999	TRUE
x		0.748271026	TRUE
x		0.748267358	TRUE
x		0.747496789	TRUE
	x	0.747299926	TRUE
	x	0.746983345	TRUE
x		0.746207248	TRUE
x		0.74530661	TRUE
x		0.5	FALSE
x		0.5	FALSE
x		0.5	FALSE

9. Model: MLP - Architecture

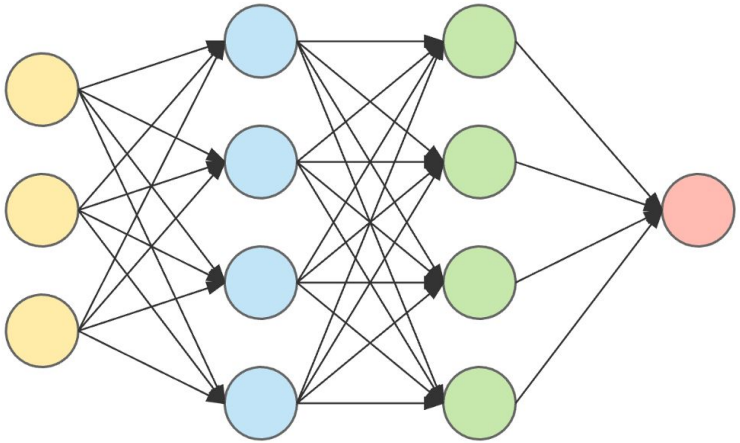

How does that make sense ???

Hidden Layers			Nodes per Hidden Layer			Results	
1	2	3	10	20	50	Balanced acc	Clears result ga
	x		x			0.752983544	TRUE
	x			x		0.752917655	TRUE
		x			x	0.752209507	TRUE
x			x			0.752151507	TRUE
x				x		0.751377622	TRUE
		x		x		0.751052197	TRUE
x					x	0.751049684	TRUE
	x				x	0.751041243	TRUE
	x				x	0.750471004	TRUE
	x			x		0.750403557	TRUE
	x		x			0.750332744	TRUE
		x		x		0.750138846	TRUE
	x				x	0.749946103	TRUE
x					x	0.749814174	TRUE
x			x			0.749810506	TRUE
x				x		0.749680387	TRUE
x				x		0.749498346	TRUE
x					x	0.749498246	TRUE
		x	x			0.749484076	TRUE
		x	x			0.749174228	TRUE
	x			x		0.749172419	TRUE
x			x			0.749171112	TRUE
		x	x			0.749039033	TRUE
	x		x			0.74885026	TRUE
	x		x			0.748849054	TRUE
	x			x		0.748721999	TRUE
		x		x		0.748271026	TRUE
x				x		0.748267358	TRUE
x					x	0.747496789	TRUE
x			x			0.747299926	TRUE
		x			x	0.746983345	TRUE
		x			x	0.746207248	TRUE
	x				x	0.74530661	TRUE
		x	x			0.5	FALSE
		x		x		0.5	FALSE
		x			x	0.5	FALSE

9. Model: MLP - Optimizers

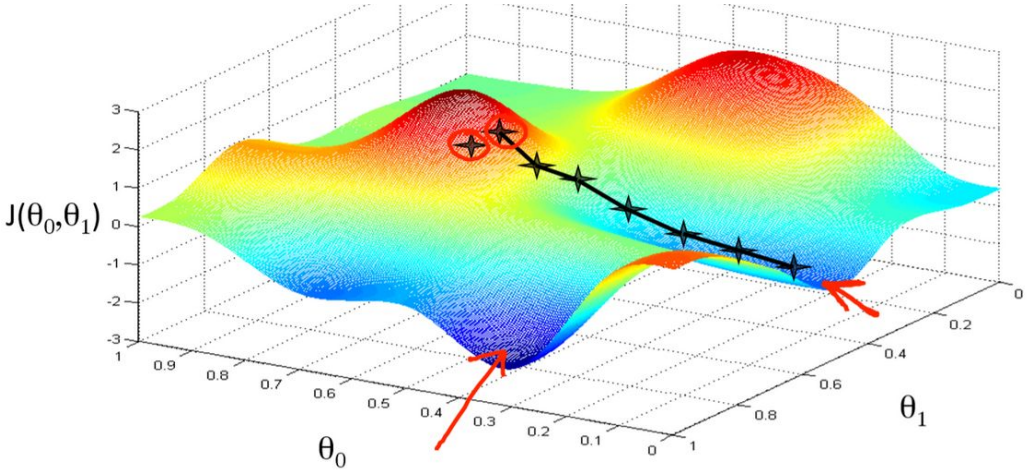

Adaptive Moment Estimation
vs
Stochastic Gradient Descent

Solver		Results	
adam	sgd	Balanced acc	Clears result g
x		0.752983544	TRUE
x		0.752917655	TRUE
	x	0.752209507	TRUE
x		0.752151507	TRUE
x		0.751377622	TRUE
x		0.751052197	TRUE
x		0.751049684	TRUE
	x	0.751041243	TRUE
x		0.750471004	TRUE
x		0.750403557	TRUE
x		0.750332744	TRUE
	x	0.750138846	TRUE
	x	0.749946103	TRUE
x		0.749814174	TRUE
x		0.749810506	TRUE
x		0.749680387	TRUE
	x	0.749498346	TRUE
	x	0.749498246	TRUE
x		0.749484076	TRUE
x		0.749174228	TRUE
	x	0.749172419	TRUE
	x	0.749171112	TRUE
	x	0.749039033	TRUE
	x	0.74885026	TRUE
	x	0.748849054	TRUE
	x	0.748721999	TRUE
x		0.748271026	TRUE
	x	0.748267358	TRUE
	x	0.747496789	TRUE
	x	0.747299926	TRUE
x		0.746983345	TRUE
x		0.746207248	TRUE
x		0.74530661	TRUE
	x	0.5	FALSE
	x	0.5	FALSE
	x	0.5	FALSE

9. Model: MLP - Performance

Accuracy = 0.742 for dataset with dropped NaN values

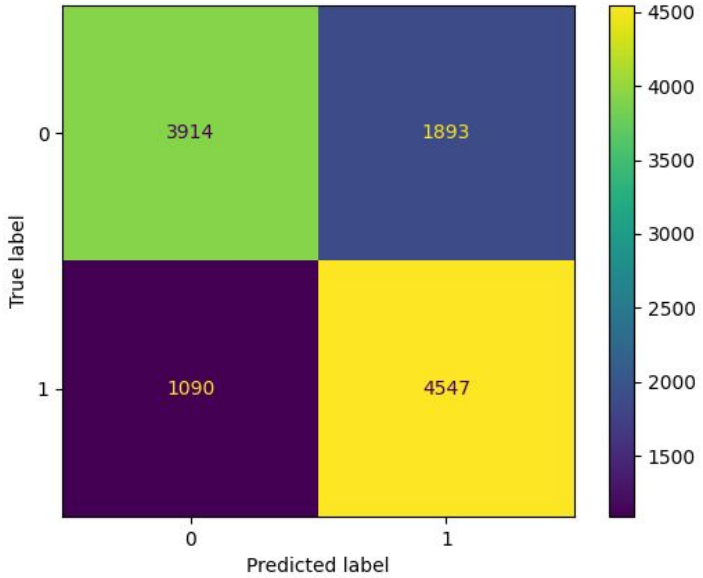

Architecture: 20, 20, 20

Accuracy = 0.753 for dataset with imputed NaN values

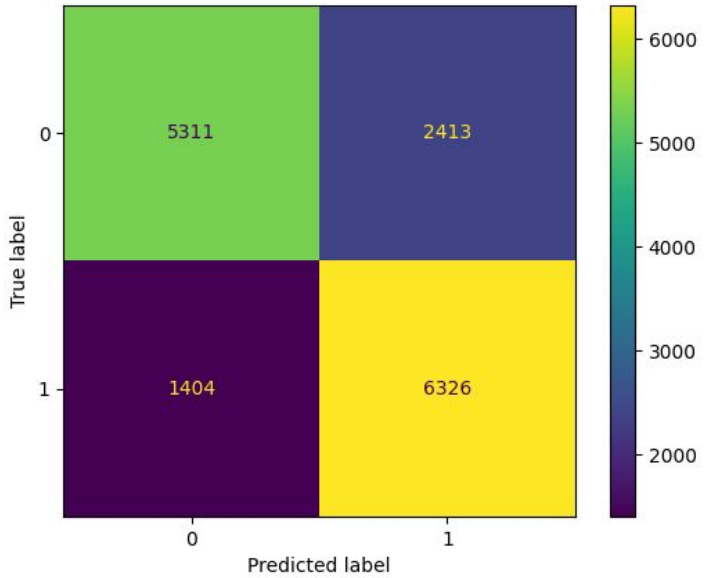

Architecture: 10, 10

10. Algorithm performance comparison

Accuracy:

- Highest balanced accuracy of implemented algorithms: 0.753 MLP(imputed)
- Worst balanced accuracy of implemented algorithms: 0.708 Decision Tree(dropped)
- Undersampling the data lead to a moderate number of false negatives

Speed:

- The optimised decision tree is very fast with predictions once trained and optimized
- KNN has the disadvantage of a huge the dataset requiring intensive computational work but is still quite fast
- MLP is extremely dependent on the hyperparameter setup

Dropped NaN vs. Imputed NaN in best case:

- Imputed NaN values consistently improved accuracy over dropped Nan

Accuracy for KNN:
Dropped = 0.72
Imputed = 0.74
Difference = 0.010

Accuracy for Decision Tree:
Dropped = 0.708
Imputed = 0.728
Difference = 0.020

Accuracy for MLP:
Dropped = 0.742
Imputed = 0.753
Difference = 0.011

11. Conclusions

- Real-world implications of data preprocessing: e.g. simplifying diabetes label could introduce hidden dangers when assessing heart disease risk
- Preprocessing of data has major influence on final result
- While using imputed values risk introducing bias, a big and representative dataset can mitigate them
- Great number of Hyperparameters can make it very hard to understand their interactions

Complex hyperparameters example

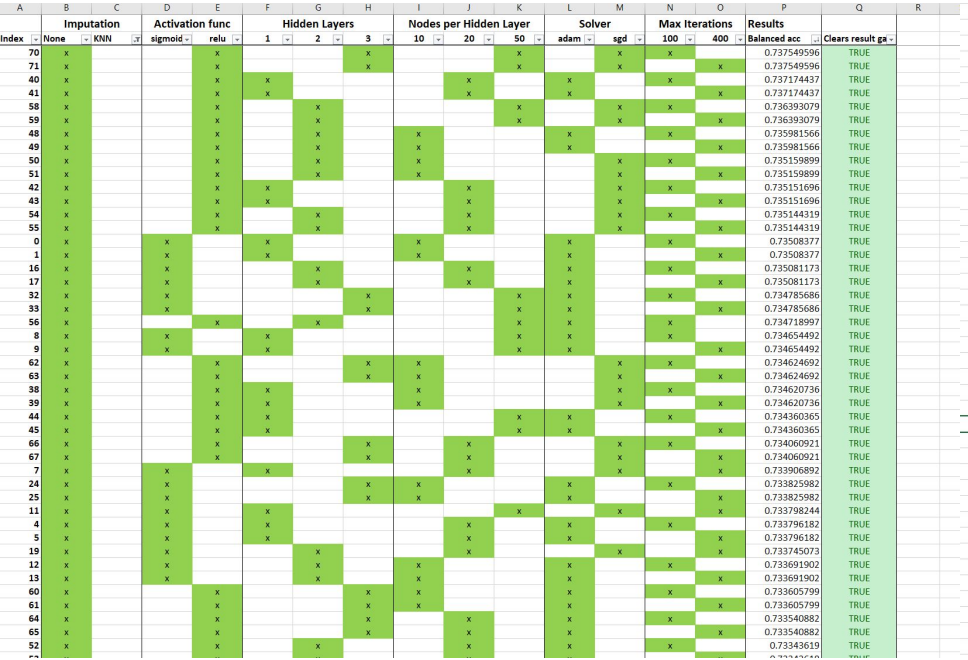

Index	Imputation		Activation func		Hidden Layers			Nodes per Hidden Layer			Solver		Max Iterations		Results	
	None	KNN	sigmoid	relu	1	2	3	10	20	50	adam	sgd	100	400	Balanced acc	Clears result ga
70	x			x			x			x		x	x		0.737549596	TRUE
71	x			x			x			x		x		x	0.737549596	TRUE
40	x			x	x				x		x		x		0.737174437	TRUE
41	x			x	x				x		x			x	0.737174437	TRUE
58	x			x		x				x		x	x		0.736393079	TRUE
59	x			x		x				x		x		x	0.736393079	TRUE
48	x			x		x		x			x		x		0.735981566	TRUE
49	x			x		x		x			x			x	0.735981566	TRUE
50	x			x		x		x				x	x		0.735159899	TRUE
51	x			x		x		x				x		x	0.735159899	TRUE
42	x			x	x				x			x	x		0.735151696	TRUE
43	x			x	x				x			x		x	0.735151696	TRUE
54	x			x		x			x			x	x		0.735144319	TRUE
55	x			x		x			x			x		x	0.735144319	TRUE
0	x		x		x			x			x		x		0.73508377	TRUE
1	x		x		x			x			x			x	0.73508377	TRUE
16	x		x			x			x		x		x		0.735081173	TRUE
17	x		x			x			x		x			x	0.735081173	TRUE
32	x		x				x			x	x		x		0.734785686	TRUE
33	x		x				x			x	x			x	0.734785686	TRUE
56	x			x		x				x	x		x		0.734718997	TRUE
8	x		x		x					x	x		x		0.734654492	TRUE
9	x		x		x					x	x			x	0.734654492	TRUE
62	x			x			x	x				x	x		0.734624692	TRUE
63	x			x			x	x				x		x	0.734624692	TRUE
38	x			x	x			x				x	x		0.734620736	TRUE
39	x			x	x			x				x		x	0.734620736	TRUE
44	x			x	x					x	x		x		0.734360365	TRUE
45	x			x	x					x	x			x	0.734360365	TRUE
66	x			x			x		x			x	x		0.734060921	TRUE
67	x			x			x		x			x		x	0.734060921	TRUE
7	x		x		x				x			x		x	0.733906892	TRUE
24	x		x				x	x			x		x		0.733825982	TRUE
25	x		x				x	x			x			x	0.733825982	TRUE
11	x		x		x					x		x		x	0.733798244	TRUE
4	x		x		x				x		x		x		0.733796182	TRUE
5	x		x		x				x		x			x	0.733796182	TRUE
19	x		x			x			x			x		x	0.733745073	TRUE
12	x		x			x		x			x		x		0.733691902	TRUE
13	x		x			x		x			x			x	0.733691902	TRUE
60	x			x			x	x			x		x		0.733605799	TRUE
61	x			x			x	x			x			x	0.733605799	TRUE
64	x			x			x		x		x		x		0.733540882	TRUE
65	x			x			x		x		x			x	0.733540882	TRUE
52	x			x		x			x		x		x		0.73343619	TRUE
53	x			x		x			x		x			x	0.73343619	TRUE
23	x				x				x			x		x	0.733304171	TRUE
68	x					x				x		x	x		0.733153686	TRUE
3	x				x		x			x		x		x	0.733071417	TRUE
57	x					x			x			x	x		0.732865454	TRUE
69	x					x				x		x	x		0.732791387	TRUE
46	x					x		x				x		x	0.732610979	TRUE
47	x					x		x				x		x	0.732610979	TRUE
36	x					x		x			x		x		0.732208202	TRUE
37	x					x		x			x			x	0.732208202	TRUE
20	x				x					x		x	x		0.731264201	TRUE
21	x				x					x		x		x	0.731264201	TRUE
28	x				x					x	x		x		0.730995093	TRUE
29	x				x					x	x			x	0.730995093	TRUE
6	x				x			x			x			x	0.730220456	TRUE
2	x				x			x				x	x		0.729251725	TRUE
10	x				x			x				x		x	0.728945729	TRUE
22	x				x					x		x		x	0.721196048	TRUE
18	x				x					x	x			x	0.715375215	TRUE
14	x				x					x		x	x		0.577537818	FALSE
15	x				x					x		x		x	0.577537818	FALSE
26	x				x					x	x		x		0.5	FALSE
27	x				x					x	x			x	0.5	FALSE
30	x				x					x		x	x		0.5	FALSE
31	x				x					x		x		x	0.5	FALSE
34	x				x					x			x		0.5	FALSE
35	x				x					x				x	0.5	FALSE
98		x			x					x	x		x		0.5	FALSE
99		x			x					x	x			x	0.5	FALSE
102		x			x					x		x	x		0.5	FALSE
103		x			x					x		x		x	0.5	FALSE
106		x			x					x			x		0.5	FALSE
107		x			x					x				x	0.5	FALSE

Questions
Feedback